

Dimensional Data Modeling

Day 1

EcZachly Inc





EcZachly Inc

What is a dimension?

- Dimensions are attributes of an entity (e.g. user's birthday, user's favorite food)
 - Some of these dimensions may IDENTIFY an entity (e.g. a user's ID)
 - Others are just attributes
- Dimensions come in two flavors
 - Slowly-changing
 - Fixed



EcZachly Inc

What we'll cover today

- Knowing your data consumer
- OLTP vs OLAP data modeling
- Cumulative Table design
- The compactness vs usability tradeoff
- Temporal cardinality explosion
- Run-length encoding compression gotchas



EcZachly Inc

Knowing your Consumer

- Data analysts / Data scientists
 - Should be very easy to query. Not many complex data types
- Other data engineers
 - Should be compact and probably harder to query. Nested types are okay
- ML models
 - Depends on the model and how its trained
- Customers
 - Should be a very easy to interpret chart



EcZachly Inc

OLTP vs master data vs OLAP

- OLTP (online transaction processing)
 - Optimizes for low-latency, low-volume queries
- OLAP (online analytical processing)
 - Optimizes for large volume, GROUP BY queries, minimizes JOINS
- Master Data
 - Optimizes for completeness of entity definitions, deduped

Mismatching needs = less business value!!!!



EcZachly Inc

- Some of the biggest problems in data engineering occur when data is modeled for the wrong consumer!



EcZachly Inc

OLTP and OLAP IS A CONTINUUM





EcZachly Inc

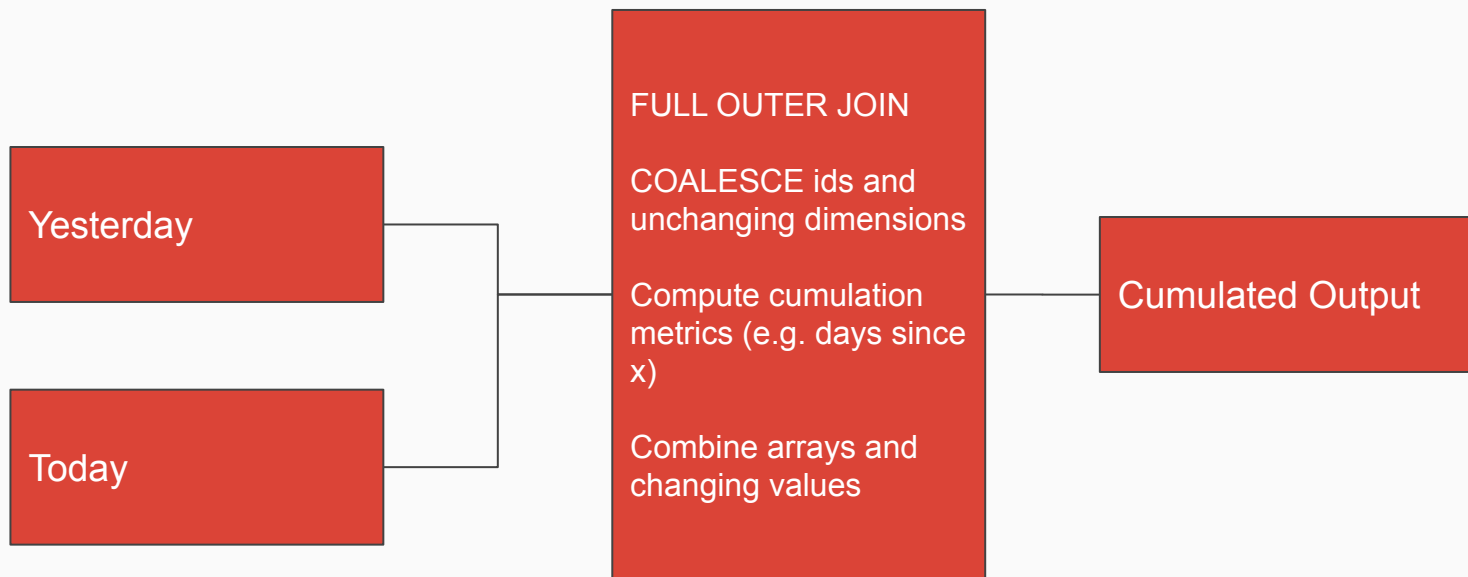
Cumulative Table Design

- Core components
 - 2 dataframes (yesterday and today)
 - FULL OUTER JOIN the two data frames together
 - COALESCE values to keep everything around
 - Hang onto all of history
- Usages
 - Growth analytics at Facebook (dim_all_users)
 - State transition tracking (**we will cover this more in Analytics track, Applying analytical patterns later**)



EcZachly Inc

Diagram of cumulative table design





EcZachly Inc

Cumulative Table Design (cont)

- Strengths
 - Historical analysis without shuffle
 - Easy “transition” analysis
- Drawbacks
 - Can only be backfilled sequentially
 - Handling PII data can be a mess since deleted/inactive users get carried forward



EcZachly Inc

The compactness vs usability tradeoff

- The most usable tables usually
 - Have no complex data types
 - Easily can be manipulated with WHERE and GROUP BY
- The most compact tables (not human readable)
 - Are compressed to be as small as possible and can't be queried directly until they're decoded
- The middle-ground tables
 - Use complex data types (e.g. ARRAY, MAP and STRUCT), making querying trickier but also compacting more



EcZachly Inc

The compactness vs usability tradeoff

- When would you use each type of table?
 - Most compact
 - Online systems where latency and data volumes matter a lot. Consumers are usually highly technical
 - Middle-ground
 - Upstream staging / master data where the majority of consumers are other data engineers
 - Most usable
 - When analytics is the main consumer and the majority of consumers are less technical



EcZachly Inc

Struct vs Array vs Map

- Struct
 - Keys are rigidly defined, compression is good!
 - Values can be any type
- Map
 - Keys are loosely defined, compression is okay!
 - Values all have to be the same type
- Array
 - Ordinal
 - List of values that all have to be the same type


Temporal Cardinality Explosions of Dimensions



EcZachly Inc


- When you add a temporal aspect to your dimensions and the cardinality increases by at least 1 order of magnitude
- Example
 - Airbnb has ~6 million listings
 - If we want to know the nightly pricing and available of each night for the next year
 - That's $365 * 6$ million or about ~2 billion nights
 - Should this dataset be:
 - Listing-level with an array of nights?
 - Listing night level with 2 billion rows?
 - If you do the sorting right, Parquet will keep these two about same size

Badness of denormalized temporal dimensions



EcZachly Inc

If you explode it out and need to join other dimensions,
Spark shuffle will ruin your compression!



EcZachly Inc

Run-length encoding compression

- Probably the most important compression technique in big data right now
 - It's why Parquet file format has become so successful
- Shuffle can ruin this. **BE CAREFUL!**
 - Shuffle happens in distributed environments when you do JOIN and GROUP BY

	season	player_name	age	height	college
1	1996	A.C. Green	33	6-9	Oregon State
2	1997	A.C. Green	34	6-9	Oregon State
3	1998	A.C. Green	35	6-9	Oregon State
4	1999	A.C. Green	36	6-9	Oregon State
5	2000	A.C. Green	37	6-9	Oregon State



Run-length encoding compression

EcZachly Inc

	season	player_name	age	height	college
1	1996	A.C. Green 5	33	6-9 5	Oregon State5
2	1997		34		
3	1998		35		
4	1999		36		
5	2000		37		



EcZachly Inc

Spark Shuffle

After a join, Spark may mix up the ordering of the rows and ruin your compression

	season	player_name	age	height	college
1	1996	A.C. Green	33	6-9	Oregon State
2	1996	Michael Jordan	34	6-6	North Carolina
3	1997	A.C. Green	34	6-9	Oregon State
4	1997	Michael Jordan	35	6-6	North Carolina
5	1998	A.C. Green 2	35	6-9 2	Oregon State 2
6	1999	A.C. Green	36	6-9	Oregon State

Let's start the workshop!



EcZachly Inc