

# Dimensional Data Modeling

Day 2

*EcZachly Inc*





*EcZachly Inc*

# What we'll cover today

- Idempotent pipelines
- Slowly-changing dimensions



*EcZachly Inc*

# Idempotent pipelines are CRITICAL

**Your pipeline produces the same results regardless of when it's ran!!!!**

**What does idempotent mean?**

- **denoting an element of a set which is unchanged in value when multiplied or otherwise operated on by itself**



*EcZachly Inc*

# Pipelines should produce the same results

- Regardless of the day you run it
- Regardless of how many times you run it
- Regardless of the hour that you run it

# Why is troubleshooting non-idempotent pipelines hard?



*EcZachly Inc*

- Silent failure!
- You only see it when you get data inconsistencies and a data analyst yells at you



# What can make a pipeline not idempotent

*EcZachly Inc*

- INSERT INTO without TRUNCATE
  - Use MERGE or INSERT OVERWRITE every time please
- Using Start\_date > without a corresponding end\_date <
- Not using a full set of partition sensors
  - (pipeline might run when there is no/partial data)
- Not using **depends\_on\_past** for cumulative pipelines



*EcZachly Inc*

# What can make a pipeline not idempotent

- Relying on the “latest” partition of a not properly modeled SCD table
  - So much pain at Facebook, DAILY DIMENSIONS AND “latest” partition is a very bad idea
  - Cumulative table design AMPLIFIES this bug
- Relying on the “latest” partition of anything else



# The pains of not having idempotent pipelines

*EcZachly Inc*

- Backfilling causes inconsistencies between the old and restated data
- Very hard to troubleshoot bugs
- Unit testing cannot replicate the production behavior
- Silent failures





# Should you model as Slowly Changing Dimensions?

*EcZachly Inc*

- Max, the creator of Airflow HATES SCD data modeling
  - [Link](#) to Max's article about why SCD's SUCK
- What are the options here?
  - Latest snapshot
  - Daily/Monthly/Yearly snapshot
  - SCD
- How slowly changing are the dimensions you're modeling?

# Why do dimensions change?



*EcZachly Inc*

- Someone decides they hate iPhone and want Android now
- Someone migrates from team dog to team cat
- Someone migrates from USA to another country
- ETC ETC ETC



# How can you model dimensions that change?

*EcZachly Inc*

- Singular snapshots
  - BE CAREFUL SINCE THESE ARE NOT IDEMPOTENT
- Daily partitioned snapshots
- SCD Types 1,2,3
-



# The types of Slowly Changing Dimensions

*EcZachly Inc*

- Type 0
  - Aren't actually slowly changing (e.g. birth date)



# The types of Slowly Changing Dimensions

*EcZachly Inc*

- Type 1
  - You only care about the latest value
  - **NEVER USE THIS TYPE BECAUSE IT MAKES YOUR PIPELINES NOT IDEMPOTENT ANYMORE**



# The types of Slowly Changing Dimensions

*EcZachly Inc*

- Type 2
  - You care about what the value was from “start\_date” to “end\_date”
  - Current values usually have either an end\_date that is:
    - NULL
    - Far into the future like 9999-12-31
  - Hard to use:
    - Since there’s more than 1 row per dimension, you need to be careful about filtering on time
- **MY FAVORITE TYPE OF SCD**
  - The only type of SCD that is purely IDEMPOTENT



# The types of Slowly Changing Dimensions

*EcZachly Inc*

- Type 3
  - You only care about “original” and “current”
  - Benefits
    - You only have 1 row per dimension
  - Drawbacks
    - You lose the history in between original and current
  - Is this idempotent?
    - Partially, which means it's not



*EcZachly Inc*

# Which types are idempotent?

- Type 0 and Type 2 are idempotent
  - Type 0 is because the values are unchanging
  - Type 2 is but you need to be careful with how you use the **start\_date** and **end\_date** syntax!
- Type 1 isn't idempotent
  - If you backfill with this dataset, you'll get the dimension as it is now, not as it was then!
- Type 3 isn't idempotent
  - If you backfill with this dataset, it's impossible to know when to pick "original" vs "current" and you'll either





*EcZachly Inc*

# SCD2 Loading

- Load the entire history in one query
  - Inefficient but nimble
  - 1 query and you're done
- Incrementally load the data after the previous SCD is generated
  - Has the same “**depends\_on\_past**” constraint
  - Efficient but cumbersome

# Let's start the lab



*EcZachly Inc*